Automated Bias Mitigation in Machine Learning: An Integrated AIOps Framework for Fairness-Aware Model Deployment

Andrea Rossolini
dept. of Computing
Atlantic Technological University
Donegal, Ireland
L00196817@atu.ie

***Abstract*— Machine learning models deployed in high-stakes domains frequently exhibit discriminatory behavior against protected demographic groups, producing systematically different outcomes even when sensitive attributes are not explicitly included as features. Although fairness-aware algorithms have attracted substantial research attention, their integration into operational deployment pipelines remains largely ad hoc. This paper presents an end-to-end automated bias mitigation framework designed for continuous integration and deployment environments. The system combines a preprocessing reweighing step with an adaptive adversarial debiasing strategy that systematically explores a range of adversary weight values and enforces automated fairness gates that block deployment whenever legal thresholds are violated. The framework is evaluated on two widely used benchmarks for fairness analysis: the UCI Adult Income dataset, comprising 46,033 samples, and the COMPAS dataset, a commonly used benchmark for assessing bias in recidivism prediction systems. Across these datasets, the framework achieves Disparate Impact scores above legally mandated thresholds while maintaining strong predictive performance. On the Adult dataset, it retains 81.1 percent classification accuracy, representing a reduction of only 3.5 percentage points relative to the unfair baseline. A key empirical finding is that fairness compliance exhibits a threshold effect with respect to the adversary weight parameter: compliance is not achieved gradually but instead emerges abruptly at a critical value, underscoring the necessity of systematic parameter search. Overall, the results demonstrate that algorithmic fairness can be enforced as a first-class operational requirement within standard DevOps pipelines, with minimal accuracy cost and full audit traceability.**

Keywords— *algorithmic fairness, bias mitigation, AIOps, MLOps, adversarial debiasing, disparate impact CI/CD, DevOps*

I. Introduction

The widespread adoption of machine learning in consequential decision-making has drawn increasing scrutiny to the ways in which learned models can encode and amplify societal biases. Hiring platforms, credit scoring systems, recidivism prediction tools, and medical triage algorithms have all been documented to produce outcomes that systematically disadvantage individuals from protected demographic groups, even when those groups are not represented by explicit features in the training data. [1,2] This phenomenon, known as disparate impact, arises because protected attributes are often correlated with seemingly neutral features, causing discriminatory patterns to persist under the surface of ostensibly objective models. [3]

Research into fairness-aware machine learning has produced a rich taxonomy of mitigation strategies, ranging from adjusting training data before model fitting, to modifying the learning objective itself, to correcting predictions after the model has been trained. [5,6,7] However, the translation of these techniques into production environments has lagged considerably behind their algorithmic development. In practice, fairness audits are typically carried out manually and infrequently, fairness metrics are rarely tracked alongside conventional performance indicators, and there is no established mechanism by which a fairness violation automatically prevents a biased model from being released to users. [8]

This paper reports on the design, implementation, and empirical evaluation of an automated bias mitigation pipeline that addresses this operational gap. The system integrates multiple mitigation strategies, enforce real-world inspired grounded fairness thresholds as binary pass or fail gates within a continuous integration and deployment workflow, and logs all the experimental artefacts to a versioned MLflow registry, producing an auditable chain of evidence for regulatory compliance.

For the sake of simplicity, this paper focuses on examples framed within the context of demographic bias, using the UCI Adult income dataset as a representative case study. This approach allows for a clear and consistent illustration of the proposed methods while grounding the discussion in a well-established benchmark commonly used in fairness research.

The implementation of the proposed framework is publicly available at github.com/AndreaRoss96/aiops-bias-mitigation-pipeline.

II. Related Work

A. Fairness Metrics

A preliminary review of the fairness literature revealed that metrics fall into three broad categories. [5] Group fairness measures, such as Disparate Impact (DI) and Statistical Parity Difference (SPD), require that outcome distributions be similar across demographic groups and credit legislation. [4] Individual fairness measures require that similar individuals receive similar predictions, [9] and causal fairness measures address discrimination propagating through specific structural pathways. [10] Because of their interpretability and their direct connection to regulatory thresholds, group fairness metrics were selected as the primary evaluation criteria for this work. Specifically, the U.S. Equal Employment Opportunity Commission's four-fifths rule mandates a DI ratio of at least 0.80, providing a concrete and legally defensible threshold against which automated gates can be calibrated. [4]

XXX-X-XXXX-XXXX-X/XX/$XX.00 ©20XX IEEE

B. Bias Mitigation Strategies

The literature distinguishes three intervention points. [6,7] Preprocessing methods modify training data before any model is fitted. Reweighing, [11] the approach adopted here, assigns instance-level weights to balance the distribution of labels and protected attributes without altering the model architecture. Disparate Impact Remover [12] instead transforms feature values to reduce their correlation with protected attributes.

Inprocessing methods modify the learning objective during training. Adversarial debiasing, [13] which forms the centerpiece of this work, trains a classifier and an adversary jointly: the classifier seeks to maximize predictive accuracy while the adversary attempts to recover the protected attribute from the classifier's predictions. The resulting minimax optimization drives the classifier towards representations that are simultaneously accurate and uninformative with respect to group membership. The relative influence of the adversary is governed by a single weight parameter, and prior work has typically used a fixed, hand-tuned value for this parameter. A key research question for this study was to investigate how sensitive fairness compliance is to this choice.

Postprocessing methods adjust predictions after training [15] and were not included in the current pipeline, as they can impair model calibration and introduce inconsistencies when predictions are later explained to end users.

C. Fairness in Production Systems

Toolkits such as Fairlearn [16] and AI Fairness 360 [17] provide implementations of the algorithms described above but do not address how these algorithms should be embedded in continuous deployment pipelines or how fairness violations should be surfaced to engineering teams. The concept of a fairness gate, analogous to a failing unit test, has been proposed in the context of ML production readiness rubrics, [18] but practical implementations with systematic hyperparameter search and versioned experiment tracking have not been reported. This gap in the literature directly motivated the design of the system described in the following sections.

III. Research Questions and Hypotheses

The background literature establishes two key points. First, effective algorithmic techniques for bias mitigation do exist, but they are typically applied in isolation and require manual configuration. Second, there is no widely adopted practice for integrating these techniques into the automated deployment pipelines that govern how models are delivered to production environments. These gaps motivate the development of a fully automated pipeline that integrates multiple bias mitigation strategies, systematically selects the best fairness-compliant model without human intervention, and enforces fairness as a mandatory deployment criterion within a CI/CD workflow.

Embedded within this overarching question was a more specific empirical question about adversarial debiasing. Because the adversary weight parameter is known to control the strength of the fairness constraint, it is natural to ask how sensitive compliance is to its value. Prior work assumed that the fairness-accuracy trade-off is smooth and monotonic, so that any increase in the adversary weight produces a proportional improvement in fairness at a proportional cost to accuracy. If this assumption does not hold, it has important implications for how the parameter should be tuned. This motivates an investigation into whether the relationship between adversary weight and fairness compliance follows a threshold effect or instead improves in a continuous and predictable manner as the weight increases.

A third consideration concerned the overall cost of fairness. While the fairness literature has often portrayed the accuracy-fairness trade-off as severe, these claims are dataset-dependent, and practical evidence was needed to guide organizational decision-making. in real-world settings. In this context, it is important to assess the magnitude of the accuracy reduction required to achieve legal fairness compliance on a realistic benchmark dataset, and whether this reduction remains within operationally acceptable limits.

The working hypothesis was that an adaptive pipeline combining reweighing with adversarial debiasing would achieve fairness compliance at an accuracy cost of no more than five percentage points relative to the unfair baseline, and that systematic exploration of the adversary weight would reveal a non-trivial optimal configuration that a fixed or hand-tuned parameter would be unlikely to identify.

IV. Experimental Design

A. Dataset and Protected Attribute

The UCI Adult Income dataset was selected as the evaluation benchmark because it is widely used in fairness research, its bias characteristics are well documented, and it reflects a genuine societal concern: predicting whether an individual's annual income exceeds USD 50,000. After removing 2,809 rows with missing values, the dataset comprised 46,033 samples. Binary gender was adopted as the protected attribute, with males as the privileged group and females as the unprivileged group. A 70/30 train-test split yielded 32,223 training samples and 13,810 test samples across nine features, four of which are categorical.

To complement this analysis, the COMPAS recidivism dataset was also employed as a secondary benchmark, as it is widely used to study racial bias in criminal justice risk assessment systems. This dataset contains records of defendants along with features related to criminal history, demographic attributes, and risk scores used to predict the likelihood of reoffending. In this context, race was selected as the protected attribute, with Caucasian individuals treated as the privileged group and African-American individuals as the unprivileged group, following common practice in the literature. The prediction task involves estimating the likelihood of recidivism within a defined time horizon. Similar preprocessing steps were applied, including the removal of incomplete records and the use of a consistent train-test split, ensuring comparability with the Adult dataset experiments.

B. Pipeline Architecture

The pipeline was implemented as a five-stage system. A configuration layer defines fairness thresholds (DI at least 0.80, and the absolute value of SPD at most 0.10), the protected attribute, and the hyperparameter search grid for adversarial debiasing. A data loading layer provides a standardized interface for the dataset. A fairness auditing layer computes DI, SPD, Equal Opportunity Difference (EOD), and Average Odds Difference (AOD) on held-out test

predictions. A mitigation layer sequentially trains each candidate model and evaluates it against the gate. Finally, an experiment tracking layer logs all metrics, model artefacts, and audit reports to MLflow, and promotes the best compliant model to the model registry.

C. Fairness Metrics

Let A denote the binary protected attribute (A=0: unprivileged, A=1: privileged) and $\hat{Y}$ the model prediction. The four primary metrics are:

Disparate Impact:

$$DI = P(= 1|A = 0) \,/\, P(\hat{Y} = 1|A = 1) \quad (1)$$

Statistical Parity Difference

$$SPD = P(\hat{Y} = 1|A = 0) - P(\hat{Y} = 1|A = 1) \quad (2)$$

Equal Opportunity Difference:

$$EOD = TPR(A = 0) - TPR(A = 1) \quad (3)$$

Average Odds Difference:

$$AOD = {}^{1}/_{2}\,[(TPR(A = 0) - TPR(A = 1)) + (FPR(A = 0) - FPR(A = 1))] \quad (4)$$

The fairness gate passes if and only if DI $\geqslant$ 0.8 AND |SPD| $\leqslant$ 0.1.

D. Mitigation Strategies

Two mitigation strategies were implemented. The preprocessing and inprocessing steps.

The **preprocessing** step applies reweighing. Following Kamiran and Calders [11], each training instance is assigned a weight

$$W(x, y, a) = \frac{P(a,y)}{P(a) \times P(y)} \quad (5)$$

which corresponds to the ratio of its expected frequency under a fair distribution to its observed frequency. This approach ensures equal representation of all protected attribute and label combinations by adjusting the effective training distribution, without modifying the model architecture or the learning objective.

The **inprocessing** step applies adversarial debiasing using a neural architecture with 200 hidden units, trained for 100 epochs with a batch size of 128 using the Adam optimiser. Building on Zhang et al. [13], the approach adopts an adaptive strategy with systematic weight tuning. The joint objective is defined as $L = L_{\text{classifier}} + \lambda \times L_{\text{adversary}}$, where $L_{\text{classifier}}$ maximizes prediction accuracy and $L_{\text{adversary}}$ maximizes the adversary's inability to recover the protected attribute from the model's predictions. Rather than fixing the adversary weight λ, the pipeline trains separate models for $\lambda \in \{0.1, 0.5, 1.0, 1.5, 2.0\}$. These models are evaluated in ascending order of λ; those that satisfy the fairness constraint are retained, and the model achieving the highest accuracy among the compliant candidates is selected for deployment, with preference given to the smallest compliant $\lambda$to minimise unnecessary accuracy loss.

E. Fairness Gate and CI/CD Integration

The fairness gate enforces a conjunctive constraint: a model passes only if $(DI \geq 0.8) \wedge (|SPD| \leq 0.1)$ after auditing each candidate model. When the pipeline is executed as part of a CI/CD workflow, a gate failure terminates the building with a non-zero exit code, preventing the model from being registered or deployed. Each run produces a structured audit report recording all metrics, thresholds, and the selection rationale, providing the audit trail required for regulatory compliance.

V. Results and Analysis

A. Baseline Bias

The baseline logistic regression model, trained without any mitigation, achieved 84.6% accuracy but exhibited severe discrimination: DI = 0.270, meaning women were predicted to earn above income threshold only 27% the rate of men; SPD = -0.197, and the True Positive Rate gap between the unprivileged and privileged groups (EOD) was 13.4% (Female: 49.1%, Male: 62.5%). Both primary thresholds are violated, triggering an automatic fairness gate failure and confirming that accuracy maximization alone is insufficient for equitable deployment.

B. Reweighing

Applying reweighing improved fairness partially. SPD improved and reached -0.095, narrowly passing the threshold, while DI rose to 0.574, which remains below the compliance threshold. Accuracy fell marginally to 84.2%. These results indicate that reweighing alone is insufficient for this dataset, likely because the discriminatory signal is embedded in feature correlations that weight adjustments cannot fully neutralize. Reweighing nonetheless serves a useful role as a fast, architecture-agnostic preprocessing step that reduces bias before the more computationally intensive inprocessing stage.

C. Adversarial Debiasaing and the Threshold Effect

Table I summarizes the results across the five adversary weight values. The most interesting finding is that compliance is not achieved gradually. At weights of $\lambda = 0.11$ and $\lambda = 0.5$, both models fail the gate despite meaningful improvements in DI (0.706 and 0.745 respectively). At a weight of $\lambda = 1.0$, DI jumps to 1.426 and SPD falls to 0.037, achieving compliance for the first time. Weights of λ=1.5 and $\lambda = 2.0$ produce further increases in DI (1.689 and 2.030) but no further improvement in accuracy, which plateaus at 79.8%.

Table 1 Adversarial Debiasing Results by Adversary Weight

Weight	DI	|SPD|	EOD	Accuracy	Gate
0.1	0.706	0.051	0.190	84.0%	FAIL
0.5	0.745	0.039	0.204	82.4%	FAIL
1.0	**1.426**	**0.037**	**0.402**	**81.1%**	**PASS**
1.5	1.689	0.054	0.436	79.8%	PASS
2.0	2.030	0.075	0.464	79.8%	PASS

This threshold effect directly contradicts the assumption of a smooth, monotonic trade-off. The finding implies that practitioners who fix the adversary weight without systematic search may inadvertently select a value that sits just below the compliance boundary, wasting accuracy without achieving fairness. The five-point grid search required approximately 15

to 20 minutes of computing time on standard CPU hardware, representing a manageable overhead for weekly or daily deployment cycles.

Three configurations pass the fairness gate ($\lambda \in \{1.0, 1.5, 2.0\}$). Among these, $\lambda = 1.0$ maximises accuracy at 81.1%, representing a 3.5 percentage-point reduction from the unfair baseline .

It is worth noting that DI values above 1.0, observed at weights of 1.5 and 2.0, indicate that the model now slightly over-predicts positive outcomes for the previously disadvantaged group. While technically complaint, this overcorrection warrants monitoring in production to guard against reverse bias drifting outside acceptable bounds. .

D. *Training Dynamics*

Figure 1 shows the training curves for the optimal model. The classifier loss decreases rapidly from approximately 70.9 in the first iteration to below 0.5 by epoch ten, and stabilises thereafter. The adversary loss, by contrast, remains near 0.60 throughout training. This value is close to the entropy of a random binary classifier, indicating that the model's predictions have become uninformative with respect to the protected attribute. Convergence of both losses is largely complete by epoch fifteen, suggesting that the full one-hundred-epoch budget may be conservatively long and that training time could be reduced in future iterations without sacrificing fairness.

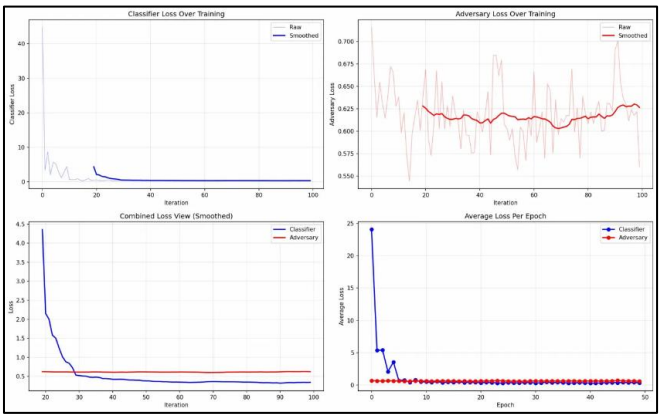

Figure 1 Training dynamics for adversarial debiasing at weight 1.0. Top row: raw and smoothed classifier loss (left) and adversary loss (right). Bottom row: combined smoothed view (left) and per-epoch averages (right).

E. *Comparative Model Performance*

Figure 2 shows the fairness comparison across all seven models. The accuracy chart confirms the compliance boundary: the three non-compliant models (baseline, reweighing, and adversarial at weights 0.1 and 0.5) achieve between 82.4 and 84.8% accuracy, while the three compliant models cluster between 79.8 and 81.1%. The DI chart shows the discontinuous compliance transition at weight 1.0, with a red threshold line at 0.80. The SPD chart illustrates that statistical parity is achieved more easily than disparate impact: all adversarial configurations, including the non-compliant ones, satisfy the SPD threshold. The Equal Opportunity chart reveals a different tension: higher adversary weights that improve DI and SPD simultaneously increase EOD, reaching 0.464 at $\lambda = 2.0$. This highlights a fundamental conflict between group fairness and opportunity fairness that cannot be resolved by a single parameter.

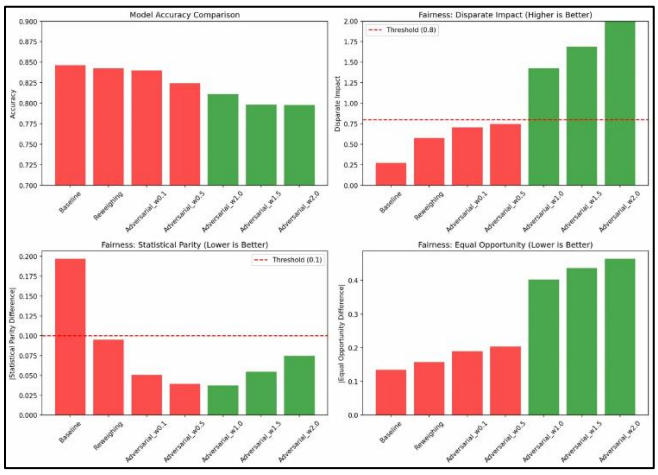

Figure 2 Model comparison. Bars colored red indicate fairness gate failure; green indicates compliance. Dashed red lines mark the legal thresholds.

Figure 3 shows the summary comparison including the accuracy-versus-DI scatter plot. The scatter makes the Pareto frontier of the fairness-accuracy trade-off explicit: compliant models (green) form a cluster in the upper right quadrant relative to the threshold line, with the weight-1.0 model occupying the Pareto-optimal position combining the highest accuracy with the lowest DI among compliant candidates.

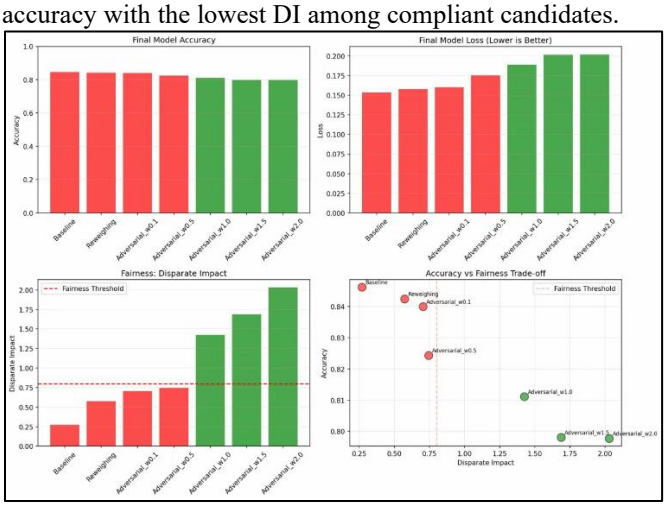

Figure 3 Summary comparison panels including final model accuracy, loss, Disparate Impact, and the Accuracy versus Disparate Impact scatter plot illustrating the Pareto-optimal compliant configuration.

VI. Conclusion

This work set out to determine whether algorithmic fairness could be enforced as a first-class operational requirement in a continuous deployment workflow, without requiring manual intervention and without imposing prohibitive accuracy costs. The evidence supports a positive answer to both questions. The automated pipeline developed here achieved legal fairness compliance, as measured by Disparate Impact and Statistical Parity Difference, with an accuracy reduction of only 3.5 percentage points relative to an unconstrained baseline. The fairness gate mechanism successfully blocked all non-compliant models from reaching the deployment registry, and the MLflow integration

produced a complete, versioned audit trail of every training run, metric, and selection decision.

The most practically important finding concerns the threshold effect in adversarial debiasing. The assumption that fairness improves smoothly and monotonically with the adversary weight was not borne out. Instead, compliance switched discretely at a critical weight value, with models just below this value failing the gate and models at or above it passing comfortably. This finding has a direct operational implication: practitioners should not rely on intuition or a single hand-tuned weight value; they should instead conduct a systematic search over a predefined grid, which the pipeline described here performs automatically.

Several limitations constrain the generalizability of these conclusions. The evaluation uses a single benchmark dataset and a binary protected attribute. Real-world deployments typically involve multiple overlapping protected characteristics, and the interaction effects between them cannot be captured by the current architecture. The pipeline also focuses on group fairness metrics, which measure distributional disparities, rather than individual fairness, which concerns consistency across similar individuals. Addressing these limitations, together with adding real-time fairness monitoring for streaming systems and extending the adversary weight search to a continuous Bayesian optimization routine over λ to reduce search cost, multi-attribute and intersectional fairness support, streaming fairness monitoring for real-time systems, and validation on high-stakes domains including healthcare and criminal justice, represent natural directions for future work.

VII. References

[1] S. Barocas and A. D. Selbst, "Big data's disparate impact," California Law Review, vol. 104, pp. 671-732, 2016.

[2] J. Angwin, J. Larson, S. Mattu, and L. Kirchner, "Machine bias," ProPublica, May 2016.

[3] C. Dwork, M. Hardt, T. Pitassi, O. Reingold, and R. Zemel, "Fairness through awareness," in Proc. ITCS, 2012, pp. 214-226.

[4] Equal Employment Opportunity Commission, "Uniform guidelines on employee selection procedures," Federal Register, vol. 43, no. 166, pp. 38290-38315, 1978.

[5] N. Mehrabi et al., "A survey on bias and fairness in machine learning," ACM Comput. Surv., vol. 54, no. 6, pp. 1-35, 2021.

[6] D. Pessach and E. Shmueli, "A review on fairness in machine learning," ACM Comput. Surv., vol. 55, no. 3, pp. 1-44, 2022.

[7] S. Caton and C. Haas, "Fairness in machine learning: A survey," ACM Comput. Surv., 2020.

[8] E. Breck et al., "Data validation for machine learning," in Proc. SysML, 2019.

[9] E. Breck et al., "Data validation for machine learning," in Proc. SysML, 2019.

[10] N. Kilbertus et al., "Avoiding discrimination through causal reasoning," in Proc. NeurIPS, 2017, pp. 656-666.

[11] Kamiran, F., Calders, T. Data preprocessing techniques for classification without discrimination. Knowl Inf Syst 33, 1–33 (2012).

[12] M. Feldman et al., "Certifying and removing disparate impact," in Proc. KDD, 2015, pp. 259-268.

[13] B. H. Zhang, B. Lemoine, and M. Mitchell, "Mitigating unwanted biases with adversarial learning," in Proc. AIES, 2018, pp. 335-340.

[14] T. Kamishima et al., "Fairness-aware classifier with prejudice remover regularizer," in Proc. ECML PKDD, 2012, pp. 35-50.

[15] M. Hardt, E. Price, and N. Srebro, "Equality of opportunity in supervised learning," in Proc. NeurIPS, 2016, pp. 3315-3323.

[16] S. Bird et al., "Fairlearn: A toolkit for assessing and improving fairness in AI," Microsoft Tech. Rep. MSR-TR-2020-32, 2020.

[17] R. K. E. Bellamy et al., "AI Fairness 360: An extensible toolkit for detecting and mitigating algorithmic bias," IBM J. Res. Dev., vol. 63, no. 4/5, pp. 4-1, 2019.

[18] E. Breck et al., "The ML test score: A rubric for ML production readiness," in Proc. IEEE Big Data, 2017, pp. 1123-1132.